

---

# Recognizing Handwritten Digits

---

SDS385

Elisa Ferracane - ef5884

Ashutosh Singh - as79592

Chakameh Jafari - cz5599

## 1 Summary

The goal of this project is to recognize handwritten digits 0-9 using the well-known MNIST dataset LeCun et al. [1998]. While very low test error rates have already been achieved (lowest is 0.23 using convolutional neural networks Ciregan et al. [2012]), this project uses only a subset of the data: 5000 for training, and 2500 for test. Because of the smaller size of the dataset, we speculated Support Vector Machine (SVM) might outperform a neural network, which is typically data-hungry. While we were able to achieve a test error rate of 0.0476 using a non-linear kernel with SVM (3rd degree polynomial), we did not have success implementing a neural network model for comparison.

## 2 SVM Models

We first removed predictors with little variability<sup>1</sup>, resulting in dropping almost 41% (319) of the predictors. Next, a grid search with 5-fold cross validation was performed over the cost values [1, 5, 10, 1000] and gamma values [0.0001, 0.001, 0.01, 1]. We experimented with linear kernel, radial, sigmoid and polynomial of degrees 3-9, with the results recorded in Table 1. Strangely, a cost of 1 was often selected by the grid search, which typically yields poor results. We obtained the best results for cost=1000, gamma=0.01 and polynomial kernel of degree 3. Finally, we created a model with these tuned parameters and tested on the test set. The resulting test error rate was **0.0476**. However, we did not conduct a statistical test to determine whether this result is statistically significantly better than the next-best model using the radial kernel with 0.0484. The R code is included in the Appendix.

Kernel	Gamma	Cost	Train error rate	Test error rate
polynomial deg=3	0.01	1	0.0592	0.22
polynomial deg=5	0.01	1	0.1284	0.22
polynomial deg=7	0.01	1	0.2612	0.22
polynomial deg=9	0.01	1	0.406	0.22
polynomial deg=9	0.01	1000	0.0498	<b>0.0476</b>
radial	0.001	5	0.0596	0.0484
sigmoid	0.001	1	0.091	0.0636

Table 1: Error rates for different SVM models, with parameters determined by grid search.

## 3 Other models

We further attempted Lasso and a feed-forward network, but encountered errors programming them in R. Finally, we took a different approach to eliminating predictors with little variability by first doing a principal components analysis and considering only the first 75 components. These were then used as input to an SVM model with a radial kernel. Although we achieved a superior test error rate of 0.02, this code was implemented in Python using the full dataset, and we were not able to replicate the results in R.

## References

Dan Ciregan, Ueli Meier, and Jürgen Schmidhuber. Multi-column deep neural networks for image classification. In *Computer Vision and Pattern Recognition (CVPR), 2012 IEEE Conference on*, pages 3642–3649. IEEE, 2012.

Yann LeCun, Léon Bottou, Yoshua Bengio, and Patrick Haffner. Gradient-based learning applied to document recognition. *Proceedings of the IEEE*, 86(11):2278–2324, 1998.

---

<sup>1</sup>SD<20, this parameter was set by tuning on the training set.

## Appendix

```
#read part of data
train_rows = 5000
test_rows  = 2500

train = read.csv("train.csv", header = F, nrow=train_rows )
test  = read.csv("test.csv", header = F, nrow=test_rows )

#convert last column to factor
train$V785 = as.factor(train$V785)
test$V785  = as.factor(test$V785)

#remove predictors with little variability
###
### 20 was sort of a guess!
###
cut_off      = 20
bad_columns = which(apply(train[,-ncol(train)], 2, sd) < cut_off)
length(bad_columns)

train = train[,-bad_columns]
test  = test[,-bad_columns]

#train and test with SVM
library(e1071)

#tune
set.seed(1010)
system.time({
  tune_svm = tune(svm , V785 ~ ., data = train , kernel ="polynomial", degree=3,
    ranges =list(cost = c(1 ,5,10,1000), gamma = c( 0.0001, 0.001, 0.01, 1 )),
    tunecontrol = tune.control(cross = 5) )
})

tune_svm

#create final model with tuned parameters
svm_final = svm(V785 ~ ., data = train , cost = tune_svm$best.parameters$cost , kernel="polynomial",
  gamma = tune_svm$best.parameters$gamma)
svm_y_final = predict(svm_final , newdata = test)

confusion_matrix_final = table(predicted = svm_y_final , actual = test$V785)
confusion_matrix_final

error_rate_final = 1 - ( sum(diag(confusion_matrix_final)) / nrow(test))
error_rate_final
```